

Assignment 2 — Due: Sept. 18 at 10 PM

tldr: Perform binary classification on the spirals dataset using a multi-layer perceptron. You must generate the data yourself.

Problem Statement Consider a set of examples with two classes and distributions as in Figure 2. Given the vector $x \in \mathbb{R}^2$ infer its target class $t \in \{0, 1\}$. As a model use a multi-layer perceptron f which returns an estimate for the conditional density $p(t = 1 | x)$:

$$f: \mathbb{R}^2 \rightarrow [0, 1] \quad (5)$$

parametrized by some set of values θ . All of the examples in the training set should be classified correctly (i.e. $p(t = 1 | x) > 0.5$ if and only if $t = 1$). Produce one plot. Show the examples and the boundary corresponding to $p(t = 1 | x) = 0.5$. The plot must be of suitable visual quality. It may be difficult to find an appropriate functional form for f , write a few sentences discussing your various attempts.

Requirements

1. Generate data using an instance of `numpy.random.Generator`. Not how many times my spirals lap the origin.
2. Create an MLP class. The MLP class should inherit from `tf.Module` and can use the `Linear` class from the previous assignment. It should have the following interface:

```
MLP(  
    num_inputs,  
    num_outputs,  
    num_hidden_layers,  
    hidden_layer_width,  
    hidden_activation=tf.identity,  
    output_activation=tf.identity,  
)
```
3. Learn how to use `sklearn.inspection.DecisionBoundaryDisplay`
4. Your network must operate on Cartesian coordinates. Do not transform the coordinates to be polar.




Figure 2: Sample spiral data.